

# **PEAK**

**P**roxy **E**liminating **A**rchitecture using **K**ubernetes

A research proposal

Joar Heimonen  
`contact@joar.me`

February 19, 2025

# Contents

<b>1</b>	<b>Part 1</b>	<b>3</b>
1.1	Introduction . . . . .	3
1.2	Background . . . . .	3
<b>2</b>	<b>Part 2</b>	<b>4</b>
2.1	Relevant Work . . . . .	4
2.2	Initial Analysis . . . . .	6
<b>3</b>	<b>Part 3</b>	<b>6</b>
3.1	Research Questions . . . . .	6
3.2	Reference Implementations . . . . .	6
	<b>References</b>	<b>7</b>

# 1 Part 1

## 1.1 Introduction

With the adoption of IPv6 [5] our networking landscape is about to change. Last year we proposed the architecture **PEAK** [4] which leverages the new capabilities of IPv6 to create distributed systems that are not reliant on proxies. Our last paper introduced a reference implementation of the PEAK architecture, while this is useful it does not allow us to benchmark and compare the performance of the PEAK architecture to traditional systems.

In this paper we propose a set of reference implementations of the PEAK architecture that will be benchmarked against a set of traditional systems. We will also propose a set of metrics that will be used to compare the performance of the systems. With the aim of showing that the PEAK architecture not only is viable but matches or exceeds the performance of traditional systems.

## 1.2 Background

Since the start of this century, the internet has grown from a small network of universities to a global network that connects billions of devices. Due to our limited address space provided by IPv4 [7] we have not been able to give every device a unique address. This has led to the development of two technologies. Network address translation (NAT) [6] and proxies.

NAT allows for the creation of subnets that can be used to connect multiple devices to the internet through a single shared IP address. This is a great solution for small networks but leads to an increase in complexity as the network grows. In a subnet consisting of servers none of these servers can be reached from the internet without the use of a proxy. The proxy simply redirects traffic to the correct in the subnet based on arbitrary rules.

While proxies works fine they introduce a single point of failure in distributed systems. If the proxy stops working, all the servers within the subnet become unreachable.

## 2 Part 2

### 2.1 Relevant Work

As mentioned in the PEAK article [4] distributed systems with DNS for load balancing and routing is not a new concept. However, it has proven to be a challenge to find recent research on the topic.

Table 1 shows the search terms we used and the results from these searches. As can be seen there is a lack of research on the topic. Our exclusion parameters were set to exclude any research older than 2020. We also excluded any research that did not mention DNS, load balancing or distributed systems in the abstract.

DB	Keywords	Date	Res.	Rel.
IEEE	Dist. Sys. AND DNS	17-Feb-25	117	0
IEEE	Load bal. AND DNS	17-Feb-25	33	0
IEEE	Dist. Sys. AND Load bal. AND DNS	17-Feb-25	13	0
Google Scholar*	CERN AND Load bal. AND DNS	17-Feb-25	40	2
ACM	Dist. Sys. AND DNS	18-Feb-25	388	0
ACM	Load bal. AND DNS	18-Feb-25	161	0
ACM	Dist. Sys. AND Load bal. AND DNS	18-Feb-25	154	0

Table 1: Search results \**Google Scholar is a meta-database, The results are from the first four pages.*

Due to the lack of search results we quickly realized that it would be beneficial to use a meta-database such as Google Scholar. This gave us two relevant articles from CERN.

- 
- **Title:** Achieving metric oriented load balancing [2]
  - **Authors:** P. Canilho, I. Reguero, P. Saiz
  - **Date:** 2020
  - **Data:** "This article presented the objectives, advantages and usage of a metric-based load balancing solution that is actively maintained at CERN." [2]
  - **Methods:** Theoretical Paper
  - **Main Findings\*:** Implemented metric-based load balancing through SNMP and their own lbclient. Does not mention how this can be applied to the larger IPv6 space.

- 
- **Title:** A Modernized Architecture for the Post Mortem System at CERN [1]
  - **Authors:** J. Barth, F. Bogyai, J. Garnier, M. Majewski, T. Martins Riberio, A. Mnich, M. Pocwierz, R. Selvek, R. Simpsons, A. Stanis, D. Wollmann, M. Zerlauth
  - **Date:** 2022
  - **Data:** Redesign of the post mortem system at the large hadron collider [1].
  - **Methods:** Theoretical Paper
  - **Main Findings\*:** Implemented load balancing through DNS to make the post mortem system of the large hadron collider horizontally scalable.

---

\*Main findings in relation to the PEAK architecture [4].

## **2.2 Initial Analysis**

We were able to find two articles that address load balancing through DNS and metric based load balancing [3]. However, none of these articles mention the use of IPv6 to create a public distributed system that is not reliant on proxies. Neither of the articles perform any form of benchmarking to compare the performance of their systems to more traditional systems.

# **3 Part 3**

## **3.1 Research Questions**

We aim to answer the following questions:

1. How does the PEAK architecture perform compared to traditional systems?
2. What metrics can be used to compare the performance of the PEAK architecture to traditional systems?
3. How does the PEAK architecture scale compared to traditional systems?

To achieve this goal this paper will propose a set of reference implementations of the PEAK architecture and a set of traditional counterparts. We will also propose a set of metrics that can be used to compare the performance of the systems.

## **3.2 Reference Implementations**

We propose the following reference implementations:

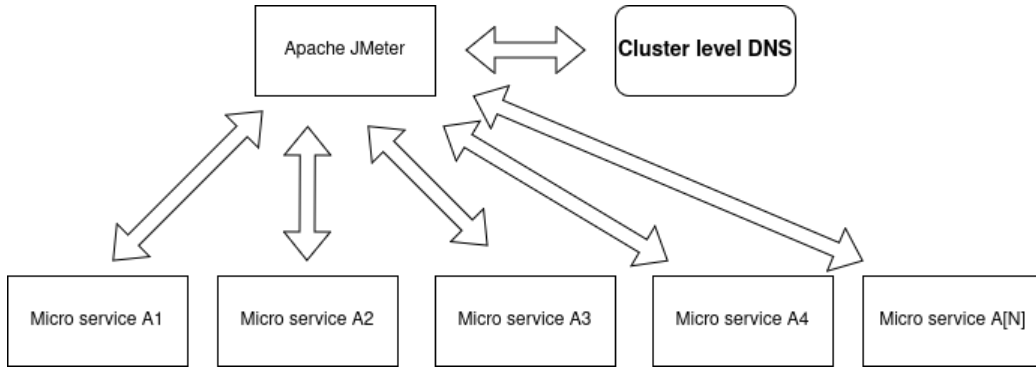


Figure 1: PEAK reference implementation 1

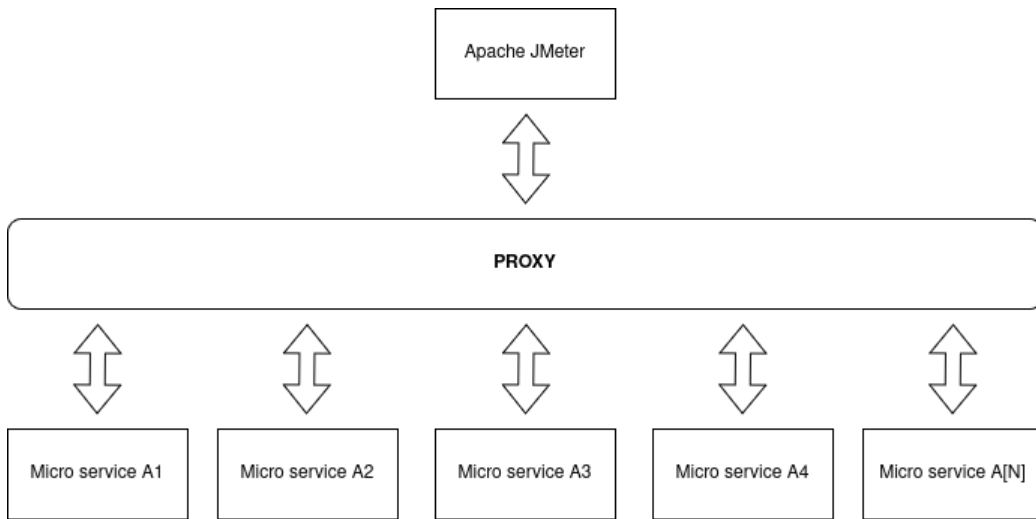


Figure 2: Reference implementation 1 traditional counterpart

## References

- [1] Jonas Barth et al. “A Modernized Architecture for the Post Mortem System at CERN”. in: (2022), 4 pages, 0.219 MB. ISSN: 2673-5490. DOI: [10.18429/JACOW-IPAC2022-TUPOMS055](https://doi.org/10.18429/JACOW-IPAC2022-TUPOMS055). (Visited on 02/19/2025).

- [2] Paulo Canilho, Ignacio Reguero, and Pablo Saiz. “Achieving Metric Oriented Load Balancing”. In: *EPJ Web of Conferences* 245 (2020). Ed. by C. Doglioni et al., p. 07004. ISSN: 2100-014X. DOI: [10.1051/epjconf/202024507004](https://doi.org/10.1051/epjconf/202024507004). (Visited on 02/19/2025).
- [3] Joar Heimonen. *PeakDNS*. Dec. 2024. (Visited on 12/11/2024).
- [4] Joar Heimonen. *Preprint PEAK Proxy Eliminating Architecture Using Kubernetes: Leveraging Proven Technologies to Create the Distributed System of the Future*.
- [5] Bob Hinden and Steve E. Deering. *Internet Protocol, Version 6 (IPv6) Specification*. Request for Comments RFC 2460. Internet Engineering Task Force, Dec. 1998. DOI: [10.17487/RFC2460](https://doi.org/10.17487/RFC2460). (Visited on 02/19/2025).
- [6] Matt Holdrege and Pyda Srisuresh. *IP Network Address Translator (NAT) Terminology and Considerations*. Request for Comments RFC 2663. Internet Engineering Task Force, Aug. 1999. DOI: [10.17487/RFC2663](https://doi.org/10.17487/RFC2663). (Visited on 02/18/2025).
- [7] *Internet Protocol*. Request for Comments RFC 791. Internet Engineering Task Force, Sept. 1981. DOI: [10.17487/RFC0791](https://doi.org/10.17487/RFC0791). (Visited on 02/18/2025).

© 2024 Joar Heimonen

This work is licensed under a [Creative Commons Attribution-Sharealike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/).